

ASSIGNMENT 10

AIM: To study and configure firewalls using IP Tables.

LO MAPPED: LO6

THEORY:

1. Firewall

A firewall is a system designed to prevent unauthorized access to or from a private network. We can implement a firewall in either hardware or software form, or a combination of both.

Generally, the firewall has two network interfaces: one for the external side of the network, one for the internal side. Its purpose is to control what traffic is allowed to traverse from one side to the other. As the most basic level, firewalls can block traffic intended for IP addresses or server ports.

TCP network traffic moves around a network in packets, which are containers that consist of a packet header—this contains control information such as source and destination addresses, and packet sequence information—and the data (also known as a payload).

While the control information in each packet helps to ensure that its associated data gets delivered properly, the elements it contains also provides firewalls a variety of ways to match packets against firewall rules.

2. Types of Firewalls:

Three basic types of network firewalls:

- **Packet filtering**, or stateless, firewalls work by inspecting individual packets in isolation. As such, they are unaware of connection state and can only allow or deny packets based on individual packet headers.
- **Stateful firewalls** can determine the connection state of packets, which makes them much more flexible than stateless firewalls. They work by collecting related packets until the connection state can be determined before any firewall rules are applied to the traffic.
- **Application firewalls** go one step further by analysing the data being transmitted, which allows network traffic to be matched against firewall rules that are specific to individual services or applications. These are also known as proxy-based firewalls.

3. Basics of IP Tables:

Iptables is a firewall, installed by default on all official Ubuntu distributions (Ubuntu, Kubuntu, Xubuntu). When you install Ubuntu, iptables is there, but it allows all traffic by default.

The rules in IP Tables are written to deal 3 different scenarios:

1. Those packets entering your machine that are destined for your machine. (INPUT)
2. Those packets leaving your machine. (OUTPUT)
3. Those packets entering your machine, but are destined for another machine and will pass through your machine (FORWARD).

In Iptables, these scenarios are referred to as INPUT, OUTPUT, and FORWARD, respectively.

Once the traffic type has been specified, three actions may be taken:

1. ACCEPT allows packets to pass through the firewall.
2. DROP ignores the packet and sends no response to the request.
3. REJECT ignores the packet, but responds to the request with a packet denied message.

COMMANDS:

1. `sudo iptables -L` (-L: List of current filters)

```
ubuntu@ubuntu:~$ sudo iptables -L
Chain INPUT (policy ACCEPT)
target     prot opt source               destination

Chain FORWARD (policy ACCEPT)
target     prot opt source               destination

Chain OUTPUT (policy ACCEPT)
target     prot opt source               destination
```

Description: We have our three default chains (INPUT, OUTPUT, and FORWARD). We also can see each chain's default policy (each chain has ACCEPT as its default policy). We also see some column headers, but we do not see any actual rules. This is because Ubuntu does not ship with a default rule set.

Basic Iptables Options

- -A - Append this rule to a rule chain. Valid chains for what we are doing are INPUT, FORWARD and OUTPUT, but we mostly deal with INPUT in this tutorial, which affects only incoming traffic.
- -p - The connection protocol used.
- --dport - The destination port(s) required for this rule. A single port may be given, or a range may be given as start: end, which will match all ports from start to end, inclusive.

- -j - Jump to the specified target. By default, iptables allows four targets:
- ACCEPT - Accept the packet and stop processing rules in this chain.
- REJECT - Reject the packet and notify the sender that we did so, and stop processing rules in this chain.
- DROP - Silently ignore the packet, and stop processing rules in this chain.
- LOG - Log the packet, and continue processing more rules in this chain. Allows the use of the -log-prefix and --log-level options.
- -i - Only match if the packet is coming in on the specified interface.
- -I - Inserts a rule. Takes two options, the chain to insert the rule into, and the rule number it should be.
- -I INPUT 5 would insert the rule into the INPUT chain and make it the 5th rule in the list.
- -v - Display more information in the output. Useful for if you have rules that look similar without using -v.
- -s --source - address[/mask] source specification
- -d --destination - address[/mask] destination specification
- -o --out-interface - output name[+] network interface name ([+] for wildcard)

Allowing Incoming Traffic on Specific Ports

We could start by blocking traffic, but you might be working over SSH, where you would need to allow SSH before blocking everything else.

To allow incoming traffic on the default SSH port (22), we could tell iptables to allow all TCP traffic on that port to come in.

2. `sudo iptables -A INPUT -p tcp --dport ssh -j ACCEPT`
`sudo iptables -L`

```

ubuntu@ubuntu:~$ sudo iptables -A INPUT -p tcp --dport ssh -j ACCEPT
ubuntu@ubuntu:~$ sudo iptables -L
Chain INPUT (policy ACCEPT)
target     prot opt source                destination            tcp dpt:ssh
ACCEPT     tcp  --  anywhere              anywhere               tcp dpt:ssh
ACCEPT     tcp  --  anywhere              anywhere               tcp dpt:ssh
ACCEPT     tcp  --  anywhere              anywhere               tcp dpt:ssh

Chain FORWARD (policy ACCEPT)
target     prot opt source                destination

Chain OUTPUT (policy ACCEPT)
target     prot opt source                destination

```

Description: Referring to the list above, we can see that this tells iptables:

- append this rule to the input chain (-A INPUT) so we look at incoming traffic
- check to see if it is TCP (-p tcp).
- if so, check to see if the input goes to the SSH port (--dport ssh).
- if so, accept the input (-j ACCEPT).

Now, let us allow all incoming web traffic

3. `sudo iptables -A INPUT -p tcp --dport 80 -j ACCEPT`
`sudo iptables -L`

```

ubuntu@ubuntu:~$ sudo iptables -A INPUT -p tcp --dport 80 -j ACCEPT
ubuntu@ubuntu:~$ sudo iptables -L
Chain INPUT (policy ACCEPT)
target     prot opt source                destination            tcp dpt:ssh
ACCEPT     tcp  --  anywhere              anywhere               tcp dpt:ssh
ACCEPT     tcp  --  anywhere              anywhere               tcp dpt:ssh
ACCEPT     tcp  --  anywhere              anywhere               tcp dpt:ssh
ACCEPT     tcp  --  anywhere              anywhere               tcp dpt:http

Chain FORWARD (policy ACCEPT)
target     prot opt source                destination

Chain OUTPUT (policy ACCEPT)
target     prot opt source                destination

```

Description: We have specifically allowed tcp traffic to the ssh and web ports, but as we have not blocked anything, all traffic can still come in.

Blocking Traffic: Once a decision is made to accept a packet, no more rules affect it.

As our rules allowing ssh and web traffic come first, if our rule to block all traffic comes after them, we can still accept the traffic we want.

All we need to do is put the rule to block all traffic at the end.

4. `sudo iptables -A INPUT -j DROP`
`sudo iptables -L`

```
ubuntu@ubuntu:~$ sudo iptables -A INPUT -j DROP
ubuntu@ubuntu:~$ sudo iptables -L
Chain INPUT (policy ACCEPT)
target     prot opt source                destination            tcp dpt:ssh
ACCEPT     tcp  --  anywhere               anywhere               tcp dpt:ssh
ACCEPT     tcp  --  anywhere               anywhere               tcp dpt:ssh
ACCEPT     tcp  --  anywhere               anywhere               tcp dpt:http
DROP       all  --  anywhere               anywhere

Chain FORWARD (policy ACCEPT)
target     prot opt source                destination

Chain OUTPUT (policy ACCEPT)
target     prot opt source                destination
```

Description: Because we did not specify an interface or a protocol, any traffic for any port on any interface is blocked, except for web and ssh.

Editing iptables: An only problem with our setup so far is that even the loopback port is blocked. We could have written the drop rule for just eth0 by specifying `-i eth0`, but we could also add a rule for the loopback. If we append this rule, it will come too late after all the traffic has been dropped. We need to insert this rule before that. Since this is a lot of traffic, we will insert it as the first rule so it is processed first.

5. `sudo iptables -I INPUT 1 -i lo -j ACCEPT` `sudo iptables -L`

```
ubuntu@ubuntu:~$ sudo iptables -I INPUT 1 -i lo -j ACCEPT
ubuntu@ubuntu:~$ sudo iptables -L
Chain INPUT (policy ACCEPT)
target     prot opt source                destination            tcp dpt:ssh
ACCEPT     all  --  anywhere               anywhere               tcp dpt:ssh
ACCEPT     tcp  --  anywhere               anywhere               tcp dpt:ssh
ACCEPT     tcp  --  anywhere               anywhere               tcp dpt:ssh
ACCEPT     tcp  --  anywhere               anywhere               tcp dpt:http
DROP       all  --  anywhere               anywhere

Chain FORWARD (policy ACCEPT)
target     prot opt source                destination

Chain OUTPUT (policy ACCEPT)
target     prot opt source                destination
```

We

will list iptables in greater detail.

6. sudo iptables -L -v

```
ubuntu@ubuntu:~$ sudo iptables -L -v
Chain INPUT (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target    prot opt in     out     source               destination
    6   464 ACCEPT    all  --  lo      any      anywhere             anywhere
    0     0 ACCEPT    all  --  lo      any      anywhere             anywhere
    0     0 ACCEPT    tcp  --  any     any      anywhere             anywhere      tcp dpt:ssh
    0     0 ACCEPT    tcp  --  any     any      anywhere             anywhere      tcp dpt:ssh
    0     0 ACCEPT    tcp  --  any     any      anywhere             anywhere      tcp dpt:ssh
    0     0 ACCEPT    tcp  --  any     any      anywhere             anywhere      tcp dpt:http
   86 181K DROP      all  --  any     any      anywhere             anywhere

Chain FORWARD (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target    prot opt in     out     source               destination

Chain OUTPUT (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target    prot opt in     out     source               destination
```

Description: You can now see a lot more information. This rule is actually very important, since many programs use the loopback interface to communicate with each other.

7. sudo iptables -A INPUT -p ICMP -j ACCEPT sudo iptables -L

```
ubuntu@ubuntu:~$ sudo iptables -A INPUT -p icmp -j ACCEPT
ubuntu@ubuntu:~$ sudo iptables -L
Chain INPUT (policy ACCEPT)
target    prot opt source               destination
ACCEPT    all  --  anywhere             anywhere
ACCEPT    all  --  anywhere             anywhere
ACCEPT    tcp  --  anywhere             anywhere      tcp dpt:ssh
ACCEPT    tcp  --  anywhere             anywhere      tcp dpt:ssh
ACCEPT    tcp  --  anywhere             anywhere      tcp dpt:ssh
ACCEPT    tcp  --  anywhere             anywhere      tcp dpt:http
DROP      all  --  anywhere             anywhere
ACCEPT    icmp --  anywhere             anywhere
ACCEPT    icmp --  anywhere             anywhere

Chain FORWARD (policy ACCEPT)
target    prot opt source               destination

Chain OUTPUT (policy ACCEPT)
target    prot opt source               destination
```

Description: Accepting ICMP Packets

8. sudo iptables -F sudo iptables -L


```

Chain INPUT (policy ACCEPT)
target     prot opt source                               destination

Chain FORWARD (policy ACCEPT)
target     prot opt source                               destination

Chain OUTPUT (policy ACCEPT)
target     prot opt source                               destination

```

Description: Clearing all rules.

9. ping 192.168.92.17

Description: You can see the response packets received. Now block incoming ICMP packets from the neighbour using command:

10. sudo iptables -A INPUT -p ICMP -j DROP
sudo iptables -L

```

ubuntu@ubuntu:~$ ping 198.168.92.17
PING 198.168.92.17 (198.168.92.17) 56(84) bytes of data.

^C
--- 198.168.92.17 ping statistics ---
22 packets transmitted, 0 received, 100% packet loss, time 21515ms

ubuntu@ubuntu:~$ sudo iptables -A INPUT -p icmp -j DROP
ubuntu@ubuntu:~$ sudo iptables -L
Chain INPUT (policy ACCEPT)
target     prot opt source                               destination
DROP      icmp -- anywhere                        anywhere

Chain FORWARD (policy ACCEPT)
target     prot opt source                               destination

Chain OUTPUT (policy ACCEPT)
target     prot opt source                               destination

```

11. Ping 192.168.92.17 sudo iptables -A INPUT -p TCP -j DROP

```

ubuntu@ubuntu:~$ ping 192.168.92.17
PING 192.168.92.17 (192.168.92.17) 56(84) bytes of data.
^C
--- 192.168.92.17 ping statistics ---
9 packets transmitted, 0 received, 100% packet loss, time 8228ms

ubuntu@ubuntu:~$ sudo iptables -A INPUT -p TCP -j DROP
ubuntu@ubuntu:~$ sudo iptables -L
Chain INPUT (policy ACCEPT)
target     prot opt source                destination
DROP       icmp -- anywhere              anywhere
DROP       tcp  -- anywhere              anywhere

Chain FORWARD (policy ACCEPT)
target     prot opt source                destination

Chain OUTPUT (policy ACCEPT)
target     prot opt source                destination
DROP       icmp -- anywhere              anywhere

```

Description: Try to ping your neighbour machine again ping 192.168.92.17. you cannot see receive ICMP echo reply packets. ping 192.168.92.17 ---does not allow (192.168.92.17 cannot send you ICMP packets) ping any other machine: ping 192.168.92.11 ---allowed (192.168.92.11 can send you ICMP packets)

Types of iptables:

I. IPTABLES TABLES and CHAINS: IP Tables has the following 4 built-in tables.

1. Filter Table: Filter is default table for iptables. So, if you do not define you own table, you will be using filter table. Iptables's filter table has the following built-in chains.

- **INPUT chain** – Incoming to firewall. For packets coming to the local server.
- **OUTPUT chain** – Outgoing from firewall. For packets generated locally and going out of the local server.
- **FORWARD chain** – Packet for another NIC on the local server.

For packets routed through the local server. Type the following command and see the result
`sudo iptables -t filter -L`

2. NAT table: Iptables's NAT table has the following built-in chains.

- **PREROUTING chain** – Alters packets before routing. i.e., Packet translation happens immediately after the packet comes to the system (and before routing). This helps to translate the destination ip address of the packets to something that matches the routing on the local server. This is used for DNAT (destination NAT).
- **POSTROUTING chain** – Alters packets after routing. i.e., Packet translation happens when the packets are leaving the system. This helps to translate the source ip

address of the packets to something that might match the routing on the destination server. This is used for SNAT (source NAT).

- OUTPUT chain – NAT for locally generated packets on the firewall. Type the following command and see the result `sudo iptables -t nat -L`

3. Mangle table: Iptables's Mangle table is for specialized packet alteration. This alters QOS bits in the TCP header. Mangle table has the following built-in chains.

- PREROUTING chain
- OUTPUT chain
- FORWARD chain
- INPUT chain
- POSTROUTING chain